
Bayesian Documentation

Release 0.1.5

Neville Newey

Oct 25, 2016

CONTENTS

1	1	Introduction	3
2	2	Statement of Problem	5
3	3	Building The Model in Python	7
3.1	3.1	Identify the Variables	7
3.2	3.2	Identify the Possible Values for the Variables	7
3.3	3.3	Create a Mental Picture of the Network	8
3.4	3.4	Create the Python Function Stubs	8
3.5	3.5	Filling in the details	9
3.6	3.6	Tieing it all Together	10
3.7	3.7	Performing Inference with our BBN	12
4		Indices and tables	15

Contents:

1 INTRODUCTION

This short tutorial will show you how to build and perform exact inference on a simple Bayesian Belief Network.

The problem we will be modelling is the famous [Monty Hall Problem](#). We choose the Monty Hall Problem as it has several features that make it ideal for a tutorial of this nature:

- It is a small problem, having only three variables
- It has an interesting and un-expected outcome (at first!)
- The conditional probabilities are simple to code in a few lines of Python

2 STATEMENT OF PROBLEM

A television game show works as follows:

A host, by the name of Monty Hall, invites a guest to select one of three doors. He explains to the guest that behind one of the doors is a car, and behind each of the other two, a goat. The host however does not tell the guest which item is hidden behind each door.

The guest gets to keep the item that is behind his choice of door.

Note: We will assume that the guest has no prior knowledge of what is hidden behind each door. Further we will assume that for most guests its more desirable to win a car than a goat!

After making his initial choice, Monty makes the game more interesting by opening one of the two remaining doors.

Note: Since Monty Hall *knows* which door conceals the car he will never open that door so he will always reveal a goat.

After opening one of the doors, Monty now offers the guest the opportunity to switch his original choice of door.

The question we are trying to answer is:

Does it matter whether the guest sticks with their original choice of door or should he switch doors?

Many people are inclined to say that it does not matter.

3 BUILDING THE MODEL IN PYTHON

Building Bayesian Belief Network models with this package is very straight forward providing we follow some simple conventions and rules.

3.1 Identify the Variables

Firstly for any problem we are trying to model it helps to identify the variables (strictly [Random Variables](#))

For the Monty Hall problem it further helps to think of the problem in terms of *events*. (This is not necessarily the case for all problems)

In the Monty hall problem, the following events occur, listed in chronological order:

1. The car is hidden behind one of the three doors prior to the guest arriving
2. The guest selects a door
3. Monty hall opens one of the two remaining doors

Each of the above events will be a **variable** in our model. Lets call these variables:

1. prize_door
2. guest_door
3. monty_door

3.2 Identify the Possible Values for the Variables

Now when we are building a Bayesian Belief Network we need to identify not only the variables of the model, but also what *values* they can take.

In this case its simple, there are three doors so lets call these **A**, **B** and **C**. Thus all three variables can take on any of the values **A**, **B** or **C**.

Note: We call the set of values a variable may take the variable's *domain*. It is not always the case that all the variables have the same domain.

3.3 Create a Mental Picture of the Network

We will now begin building the actual Bayesian Belief Network (BBN). A BBN is a Directed Acyclic Graph (DAG), which means that we need to construct a graph consisting of nodes (also called *vertices*) and directed edges.

Each variable identified in the previous section, will be a node in the network. We start off by placing a node representing each variable into our graph:




Fig. 3.1: Monty Hall Variables

Now we need to add the edges. Edges in Probabilistic Graphical Models represent the fact that two nodes connected by each edge *influence* one another. In a BBN, since the edges are directed, implying a “parent”/“child” relationship, an edge represents the fact that the child variable is conditionally dependent on the parent variable. This means that the value the child node takes, is dependent on the value (or value) of its parent (or parents). Thinking again about the events in the Monty Hall problem, when Monty opens the door to reveal the goat, his choice *depends* on which door the prize is hidden behind. (Since he will never choose the door concealing the car). Similarly Monty will never choose the same door as the guest, this would defeat the purpose of the game. We thus create an edge starting from the prize door node, and ending at the monty_door node. Similarly, we create an edge starting at the guest_door and ending at monty_door. The final structure of our BBN graph looks like this:

3.4 Create the Python Function Stubs

Now that we have constructed our graph, we are finally ready to start encoding the model in Python. In this package, every node in the graph is represented by a Python function. (Any callable will do).

The parameters of the function represent the variables in the model. For the prize_door node we thus create the following stub:

```
def f_prize_door(prize_door):  
    pass
```




Fig. 3.2: Monty Hall Nodes and Edges

Similarly for the `guest_door` node, we will create the following stub:

```
def f_guest_door(guest_door):  
    pass
```

Now for the third and final node in our graph, `monty_door`. The Python function for this node is a little more interesting. Since this node has edges coming into it from both the `prize_door` node and the `guest_door` node, we will add those two variables into the functions parameter list:

```
def f_monty_door(prize_door, guest_door, monty_door):  
    pass
```

There are several important points to note about the three Python function stubs we have introduced here:

1. The parameter names are just the same as the variable names for our model
2. The functions represent the nodes of the graph
3. To distinguish between node names and variable names we have prefixed the function names with “**f_**” (Note that this is simply a convention, you can call them anything but as you will see this convention is used in all the examples in this package.)
4. Nodes that have parents need to include the parent variables in their parameters.
5. There are exactly three functions, and exactly three variables representing the three events in the model.
6. Each function introduces one new variable into the model

3.5 Filling in the details

No BBN is complete without having both the graph *and* the distributions governing the variables. We will now start filling out the prior and conditional probabilities for each of the three variables. While doing this it helps to think of the functions in this way:

Each function must return a real number between 0 and 1, representing the probability that the variable it represents takes on a particular value in the variable's domain.

In the case of `f_prize_door`, we want to return for each possible value of **A**, **B**, **C** (ie each door) the probability that the prize is hidden behind that door. Assuming that the door is hidden randomly behind any of the three doors the probability of it being behind any particular one is $1/3$. Lets modify the stub to reflect this:

```
def f_prize_door(prize_door):  
    return 0.33333333
```

Since the guest has no knowledge of where the prize is hidden it is obvious that the `f_guest_door` function will also simply return $1/3$:

```
def f_guest_door(guest_door):  
    return 0.33333333
```

For the more interesting function, `f_monty_door`, we need to encode the likelihood of the variable **monty_door** having each of the three values **A**, **B**, **C**, **given** the value of its parent nodes **prize_door** and **guest_door**. If the guest happens to have chosen correctly, then Monty can essentially choose any of the remaining doors at random. On the other hand if the guest got it wrong, then since Monty can open neither the guests choice nor the prize door, he has only one choice: the door concealing the goat which the guest did not choose. The encoding in Python is below:

```
def f_monty_door(prize_door, guest_door, monty_door):  
    if prize_door == guest_door: # Guest was correct!  
        if prize_door == monty_door:  
            return 0 # Monty never reveals the prize  
        else:  
            return 0.5 # Monty can choose either goat door  
    elif prize_door == monty_door:  
        return 0 # Again, Monty wont reveal the prize  
    elif guest_door == monty_door:  
        return 0 # Monty will never choose the guest door  
    else:  
        # This covers all cases where  
        # the guest has *not* guessed  
        # correctly and Monty chooses  
        # the only remaining door that  
        # wont reveal the prize.  
        return 1
```

Note: There are several other variations of encoding the same function.

3.6 Tying it all Together

We will now complete the program and perform some inference on the graph. To create the actual graph, we need the function `build_bbn` from the `bbn` module. Add this line to the top of your file:

```
from bayesian.bbn import build_bbn
```

Now lets fill in the rest of the program, add this to the bottom of the file:

```
if __name__ == '__main__':  
  
    g = build_bbn(  

```

```

f_prize_door,
f_guest_door,
f_monty_door,
domains=dict(
    prize_door=['A', 'B', 'C'],
    guest_door=['A', 'B', 'C'],
    monty_door=['A', 'B', 'C']))

```

What the above piece of code does is create an instance of the BBN class. The factory function, `build_bbn` takes as parameters, all the functions representing the nodes in the graph, and an optional domains dictionary which specifies what the domain of each variable is. Note that the structure of the graph is inferred from the functions and the parameters.

The entire program should look like this:

```

from bayesian.bbn import build_bbn

def f_prize_door(prize_door):
    return 0.33333333

def f_guest_door(guest_door):
    return 0.33333333

def f_monty_door(prize_door, guest_door, monty_door):
    if prize_door == guest_door: # Guest was correct!
        if prize_door == monty_door:
            return 0 # Monty never reveals the prize
        else:
            return 0.5 # Monty can choose either goat door
    elif prize_door == monty_door:
        return 0 # Again, Monty wont reveal the prize
    elif guest_door == monty_door:
        return 0 # Monty will never choose the guest door
    else:
        # This covers all case where
        # the guest has *not* guessed
        # correctly and Monty chooses
        # the only remaining door that
        # wont reveal the prize.
        return 1

if __name__ == '__main__':

    g = build_bbn(
        f_prize_door,
        f_guest_door,
        f_monty_door,
        domains=dict(
            prize_door=['A', 'B', 'C'],
            guest_door=['A', 'B', 'C'],
            monty_door=['A', 'B', 'C']))

```

Save the above code in a file called `monty_hall.py`

3.7 Performing Inference with our BBN

Assuming that the Bayesian package has been installed correctly we will now *query* the BBN.

Run the following command:

```
$ python -i monty_hall.py
```

The `-i` command line argument causes the Python interpreter to go into interactive mode after executing all of the code in the file. This is a handy way for us to query the model. The BBN class is primarily queried through the `query` method. There is a user friendly wrapper around the `query` method called `q`. Lets start off by calling `q` without any arguments, you should see something like this:

```
>>> g.q()
+-----+-----+-----+
| Node      | Value | Marginal |
+-----+-----+-----+
| guest_door | A      | 0.333333 |
| guest_door | B      | 0.333333 |
| guest_door | C      | 0.333333 |
| monty_door | A      | 0.333333 |
| monty_door | B      | 0.333333 |
| monty_door | C      | 0.333333 |
| prize_door | A      | 0.333333 |
| prize_door | B      | 0.333333 |
| prize_door | C      | 0.333333 |
+-----+-----+-----+
>>>
```

How do we interpret this output? The `q` method essentially calls the `query` method with the same arguments it was supplied. It then formats the results from the `query` method in a nice human readable table. The columns of the table are `Node`, `Value`, and `Marginal`. (In this section we will use the terms *node* and *variable* interchangeably since every variable is represented by exactly one node) You will notice that for each variable, and for each value in that variables domain, a marginal is shown in the table. All marginals in this query have the same value of 0.33333. This is because we did not provide any *evidence* to the model and in the absence of any evidence, these marginals would indeed approximately 1 / 3 for each door.

Note: In BBN terminology we call the assignment of a value to a variable *evidence*. We call any set of assignments of zero or more variables to a value a *configuration* of the graph.

Now lets provide some evidence to the BBN and query it again. Suppose we have observed that the guest chose door A. Type the following in the Python interpreter interactive session:

```
>>> g.q(guest_door='A')
+-----+-----+-----+
| Node      | Value | Marginal |
+-----+-----+-----+
| guest_door | B      | 0.000000 |
| guest_door | C      | 0.000000 |
| guest_door* | A*     | 1.000000 |
| monty_door | A      | 0.000000 |
| monty_door | B      | 0.500000 |
| monty_door | C      | 0.500000 |
| prize_door | A      | 0.333333 |
| prize_door | B      | 0.333333 |
+-----+-----+-----+
```



```
| prize_door | C      | 0.333333 |
+-----+-----+-----+
```

Notice the changes in the marginal values *after* we have supplied the observation that the guest chose door **A**. The marginal for the variable `guest_door` having the value **A** is now 1, i.e. certainty, since that's what we observed. Likewise the marginal for `guest_door` having any of the other values **B** or **C** is now zero since we know the guest did not choose any of those doors. Notice also that the evidenced variable, and its value are marked in the table with an asterisk, reminding us of what evidence was supplied to this query. Look at the `monty_door` variable next. Notice that the marginal for the `monty_door` variable to have the value **A** is 0. This is because the rules of the game that we have encoded do not allow Monty to pick the same door as the guest. The marginals for the `monty_door` variable having the values **B** or **C** are each 0.5 reflecting the fact that without any further evidence it's equally likely for Monty to choose any of the two remaining doors.

Now let's suppose that we observe Monty opening door **B** and making the offer of a switch to the guest. Let's add this evidence to our query:

```
>>> g.q(guest_door='A', monty_door='B')
+-----+-----+-----+
| Node      | Value | Marginal |
+-----+-----+-----+
| guest_door | B     | 0.000000 |
| guest_door | C     | 0.000000 |
| guest_door*| A*    | 1.000000 |
| monty_door | A     | 0.000000 |
| monty_door | C     | 0.000000 |
| monty_door*| B*    | 1.000000 |
| prize_door | A     | 0.333333 |
| prize_door | B     | 0.000000 |
| prize_door | C     | 0.666667 |
+-----+-----+-----+
>>>
```

Notice once again that all evidenced variables and the values they have been assigned have been marked with an asterisk. Now let's look at the marginal column for the variable `prize_door`. Herein lies the answer to our original question of whether it makes a difference or not if the guest switches their choice.

As we can see, given the evidence we supplied, the prize is *twice* as likely (.667/.333) to be behind door **C** than behind door **A**. The guest should thus switch doors.

This concludes this short introductory tutorial. Future tutorials will show other functionality in this package.

Good luck and have fun building your own models!

INDICES AND TABLES

- `genindex`
- `modindex`
- `search`